



Edward Q. Earley

Engineering Leader, now building hands-on — Seattle & Chicago

WHO I AM

An engineering career with one through-line: building the platform infrastructure that removes the gatekeeping layer — taking what was reserved for a few and making it available to everyone at scale. At WMS, I created a \$30M social casino category that didn't previously exist. At SpotHero, I rebuilt infrastructure to 100x delivery capacity and 10,000x transactional throughput. At Otus, enterprise-grade analytics for K-12 schools that couldn't otherwise afford them. At AWS, I owned the tooling — EKS Addons, eksctl, ACK — that made Kubernetes accessible to teams without dedicated platform engineering. At Meta, I led the capacity-governance platform for ad-serving infrastructure and rebuilt its operations around agents and automation — then moved back to hands-on building as a Staff Production Engineer on the machine-learning infrastructure behind that same ad-serving fleet. I've owned the full stack — product direction, technical architecture, and the org that executes — and now I'm back in the serving path day to day. AI-native engineering is a first-class practice. Small teams, bold bets, first to market.

SKILL SETS

Technical

Java/JVM · Kubernetes · AWS · Terraform · Kafka · cloud-scale infrastructure, services & operations · physically and virtually distributed systems · SRE

Product & Strategy

Full-stack product + engineering leadership · category creation · roadmap ownership · competitive positioning · GTM strategy

Executive & Operations

Member of the executive team (Otus, SpotHero, WMS) · org design & transformation · P&L and budget ownership · recruited nationally recognized senior talent · built orgs from 5→100+ · exec stakeholder management

AI & Transformation

AI-native engineering practices · AI agents in infrastructure · AI-assisted IaC · AI-driven management workflows

EXPERIENCE

Meta

August 2025 – Present

Staff Production Engineer — machine-learning infrastructure behind ad serving

June 2026 – Present

A deliberate move back to hands-on building after leading the capacity-governance platform for the same ad-serving fleet. Production engineering on the machine-learning infrastructure behind Meta's ad-serving systems — reliability, performance, and the platform foundations that let ML teams ship faster.

- Focused on the reliability and performance of large-scale ML serving infrastructure — debugging latency and regressions in the serving path, and tuning the trade-offs that keep ad-serving models fast, available, and cost-efficient under continuous load.
- Applying agents and automation to production operations — compressing the loop from signal to diagnosis to fix.
- Bringing a platform-owner's lens to IC work — building for self-service and operability so reliability makes teams faster instead of taxing them.

Focus: reliability, performance, and platform foundations for ML-driven ad serving

Production Engineering Manager (Platform & AI Transformation Lead, 16-person org)

August 2025 – June 2026

Led a 16-engineer team building the capacity governance platform for Meta's ad-serving infrastructure — the system that controls how resources are allocated, enforced, and

reclaimed across dozens of services. The role spanned platform engineering, people leadership, and AI-native transformation of both our infrastructure operations and how we ran the team itself.

- Designed and shipped automation that replaced manual resource management with a self-service platform — provisioning, quota enforcement, cost attribution, and anomaly detection at a scale that supports billions in annual ad revenue.
- Introduced AI agents into infrastructure operations — systems that detect cost anomalies and recommend optimization actions, replacing manual engineering analysis with automated workflows.
- Drove an AI-native transformation of engineering management itself — using AI to automate program tracking, status reporting, team communications, stakeholder management, and decision support across a 16-person organization.
- Drove tens of millions in annual infrastructure cost savings through platform automation and AI-accelerated optimization.
- Built and organized three engineering pods from an inherited flat structure, each with clear ownership: resource optimization, fleet observability, and enforcement.
- Grew the team from 10 to 16 engineers in two quarters through active sourcing, sell calls, and onboarding. Managed the full people lifecycle — hiring, performance calibration, growth planning, and career development for engineers spanning junior through staff level.
- Full-stack product + engineering ownership — owning problem framing, roadmap trade-offs, and execution priorities directly.

Team Size: 16 engineers

Scale: ~\$30M+ in annual infrastructure cost savings

[Amazon / AWS](#)

May 2022 – August 2025

[EKS](#), Engineering Manager (End-to-end owner: EKS Addons, eksctl, ACK, KRO)

April 2024 – August 2025

Owned the core components defining how customers extend and operate EKS at scale — drove AI-native infrastructure-as-code strategy bringing declarative AWS resource management natively into Kubernetes.

- [ACK](#) & [KRO](#): declarative, Kubernetes-native management of AWS resources — enabling customers to manage RDS, S3, DynamoDB, and more via Kubernetes manifests and GitOps workflows.

- Led AI-native IaC product development: next-generation tooling bringing declarative resource management and optimization natively into the EKS operational model.

Team Size: 12+ engineers

Scale: Operated within EKS (est. ARR ~\$1B+) — the most widely adopted managed Kubernetes platform globally; owned primary customer-facing tooling surface

[Amazon Managed Grafana](#), Engineering Manager (End-to-end owner)

May 2022 – April 2024

Single-threaded owner of AWS's managed open-source observability platform — owned the full engineering lifecycle from execution to roadmap delivery.

- Delivered Custom Plugins support (re:Invent 2023) and a new ECS/Fargate dataplane (re:Invent 2022). Two consecutive flagship launches on AWS's largest engineering stage.
- Established full-lifecycle program and project mechanisms (ideation → scrum practices → health reporting → roadmap review) adopted as the standard across the broader AWS observability org.
- Built a net-new S3 line-of-business team in Austin from scratch — 1 to 9 engineers in one quarter.

Team Size: 12+ engineers

Scale: Estimated ARR ~\$50M+; end-to-end ownership of a standalone AWS service; two consecutive re:Invent flagship launches

[Wayfair](#)

Associate Director, Platform Engineering

July 2021 – April 2022

Led a ~30-person org (3 managers, 6+ senior ICs) across JVM platform, big data, DevOps tooling, and codebase management.

- Rebooted hiring by making engineering managers directly accountable for outcomes; onboarded 7 engineers in <2 months across 2 continents.
- Built the JVM Platform team from 1 engineer to a 10-person EU/US-distributed org supporting 300+ Tier 1 and Tier 2 services.
- Stood up a 3-person Apache Flink team for event-based services; shipped into production across multiple Tier 1 services.
- Introduced OKR reporting and MBRs with delivery metrics used at exec level; format replicated at parent org.

- Created a Single-Threaded Ownership program to develop future leaders; 2 engineers fast-tracked for promotion, model adopted at parent org.

Team Size: ~30 reports (3 managers, 6+ senior ICs)

Scale: Wayfair revenue ~\$14B, market cap ~\$30B (2021)

Amazon / AWS

S3, Engineering Manager (Control Plane Innovation — Ignite)

July 2020 – July 2021

High-performance incubator team within S3's Control Plane — 4 concurrent workstreams at planetary scale.

- Citadel: shipped S3's first worldwide real-time hardware failure reporting engine — baseline technology for ML and predictive analytics on global S3 fleets.
- Mechanic: tiger team that eliminated the need for SSH access to production; reduced SSH by 90%+ in four months.
- ONAC: infra-as-code migration of S3 to native AWS networking — the foundational network architecture and security model for S3.

Team Size: 5 engineers

Scale: Operated within S3 (est. ARR ~\$10B+) — stabilized planetary-scale Control Plane services; 90%+ SSH elimination

Otus

Chief Architect & Chief Data Officer

November 2018 – July 2020

Member of the executive team. Hybrid IC-executive role leading a full-company transformation to data and platform technologies at a Chicago edtech company.

- Delivered the first-generation Data Analytics Platform on time and on budget with a brand-new team in 5 months.
- Restructured all of Product Development into cross-functional squads in the first weeks; doubled engineering headcount in 6 weeks without a recruiter.
- Reset Agile and QA practices and put modern DevOps tooling and monitoring in place — all within the first 2 months.

- As IC: built a GitOps-based Kubernetes platform with CI/CD (Codefresh, Flux, EKS) and migrated environments to new AWS accounts via Terraform.
- Led Otus to [12th Best Small Company to Work For in Chicago \(Built In Chicago, 2019\)](#).

Team Size: 25+ reports — member of the executive team, co-owned engineering with CTO across QA, engineering, and program management

Scale: Transformed a seed-stage engineering org into a product development machine in under 6 months

SpotHero

VP Engineering

December 2016 – November 2018

Member of the executive team. Two-year organizational transformation of a B2C parking marketplace — rebuilt engineering from the ground up.

- Owned a 2-year transformation plan across process, people, architecture, and culture — delivered on time and on budget.
- Scaled platform from an early-stage marketplace to enterprise-ready infrastructure supporting national expansion: 100x increase in technical delivery capacity, 10,000x growth in transactional throughput.
- Defined and built a modern platform (Java, Kubernetes, Kafka); stood up a net-new Shared Services division — Data, Search, Developer Tools, and APIs teams.
- Introduced a pure managed-services DevOps model; deliberately avoided the anti-pattern of a centralized DevOps team by empowering product squads to own their infrastructure (Conway's Law).
- Rebooted QA; introduced lightweight Agile with normalized delivery metrics; served as engineering lead on a tuck-in acquisition (2017).
- Doubled engineering headcount without a dedicated recruiter; built management tier and senior IC layer; created and hit first engineering budget.

Team Size: 75+ — all engineering functions, top to bottom

Scale: Platform transformation that enabled national expansion; SpotHero acquired by Uber (valued ~\$500M at Series D)

Uptake

Director, Engineering

July 2016 – December 2016

Technical director across Data Engineering, API Gateway, Security Engineering, DevOps, and Streaming Data at a Chicago industrial IoT startup.

Team Size: 35+ staff

Scale: \$2B-valued industrial IoT unicorn (Caterpillar-backed)

WMS Industries / Scientific Games

Multiple Roles up to CTO

January 2009 – June 2016

Member of the executive team. Walked in as a software engineer; walked out as CTO — 7-year progression through Platform Engineering to executive leadership.

- Built award-winning platforms (Players Life, Play4Fun Network/SG Universe) that created the US B2B social casino category; platform grew to ~\$30M annual revenue.
- Architected a cloud-native AWS platform in Python as a low-end B2B disruptor — defined a product category that didn't exist, built with a 5-person team.
- Solely designed and implemented a Java web services platform connected to 1,000+ slot machines — with a fully automated integration test framework; architecture predicted features later found in Dropwizard and Spring Boot.
- Re-architected a heavily distributed, debt-laden platform using PaaS, microservices, and the Strangler pattern; restructured the org to mirror the architecture (Conway's Law).
- Executed a hard reboot of a broken global technology organization: doubled revenue and halved operating costs in year one through architectural reform and DevOps/CI/CD transformation.

Team Size: 100+ worldwide

Scale: WMS Industries acquired by Scientific Games for ~\$1.5B (2013); Scientific Games market cap ~\$7B at peak

Earlier roles not listed here — full employment history on [LinkedIn](#).

EDUCATION

University of Chicago

M.A., Social Sciences — [Thesis: "Mexican Modernities"](#) · 1998

University of Illinois at Chicago

B.A., Latin American Studies, cum laude · 1993

School of the Art Institute of Chicago

Fine Arts · 1987–1989

DePaul University

M.S., Computer Science [Partial] · 2002–2004

PATENTS (GRANTED)

- [Player initiated multi-player games — US 9,412,236 B2](#)
- [Jurisdictional control in a wagering game system plugin architecture — US 8,662,992 B2](#)
- [Messaging to and from wagering game machines — US 9,251,648 B2](#)
- [Mobile applications and wagering game machines — US 9,235,953 B2](#)
- [Layout elements as rendering placeholders for native wagering game applications — US 9,293,008 B2](#)

AWARDS

- 2011 — 1st Place, Player's Life, Casino Enterprise Management's Annual "Slot Floor Technology Awards"
- 2011 — Platinum Award, Player's Life, Casino Journal's "Top 20 Most Innovative Gaming Technology Products"
- 2011 — American Business Awards "Stevie", New Product or Service of the Year — Media & Entertainment
- 2013 — Best Consumer Service Technology, "Play4Fun Network," Global Gaming Business Awards
- 2014 — Top 20 Most Innovative Gaming Technology Products, "Play4Fun Network," Casino Journal
- 2015 — eGR North America Awards Freeplay Gaming Supplier, SG Universe "Play4Fun Network"
- 2019 — Otus ranked 12th [Best Small Company to Work For in Chicago](#)

LANGUAGES & INTERESTS

Spanish, some French, some Ancient Greek

Music, cinema, literature, social sciences, humanities, art, running, biking, yoga, travel

